

solution

November 28, 2024

1 MEGA UE5: Machine Learning

In this exercise, you will use Python to explore and get insights into a large cohort of COVID patient cases using traditional machine learning approaches for stratification (i.e., dimensionality reduction and clustering). You will work with EHR data (`data_ehr.csv`) that you can find in the UE5- data.zip file. There is also a `data_description.csv` that includes a description of the variables included in this EHR data set.

For the submission, you need to provide two deliverables: 1. A `.pdf` with the answers to the questions (a-f) below. 2. Your `.py` script(s).

1.1 Part 1: Data Preprocessing

Read the EHR data (`data_ehr.csv`), for instance with the Pandas library. The dataset contains missing values that are encoded with NA or empty values. Clean the dataset by removing columns where more than 30% of the values are NA.

a- How many dimensions (attributes) and instances (cases or number of patients) does the clean dataset contain? How many dimensions and instances did it contain before the cleaning?

b- What type of data does the clean dataset contain?

```
[4]: import part1 as preprocessing

data = preprocessing.read_csv('UE5-data/data_ehr.csv')
rows, columns = data.shape
print('Rows:', rows, 'Columns:', columns)

clean_data = preprocessing.clean_data(data)
clean_rows, clean_columns = clean_data.shape
print('Rows_cleaned:', clean_rows, 'Columns_cleaned:', clean_columns)
```

Rows: 1384 Columns: 130

Rows_cleaned: 1384 Columns_cleaned: 118

The file `data_ehr.csv` contains 130 dimensions (columns) and 1384 instances (rows). After preprocessing and removing columns/dimensions where more than 30% of the values are NA it only contains 118 dimensions (columns). The number of instances stays the same since removal only concerned columns.

```
[5]: print(clean_data.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1384 entries, 0 to 1383
Columns: 118 entries, to_patient_id to therapeutic.heparin.Boolean
dtypes: bool(3), float64(32), int64(1), object(82)
memory usage: 1.2+ MB
None
```

In the preprocessed “cleaned” data there are 118 columns left. Each of these has a specific data type: 3 columns-bool, 32 columns-float64, 1 columns-int64, 82 columns-object

1.2 Part 2: Clustering

Read the documentation of the scikit-learn library for clustering (<https://scikit-learn.org/1.5/modules/clustering.html>).

Try a couple of different clustering algorithms to cluster the cohort of patients, based on numerical features only. Then, decide which algorithm and with what parameters is best to use in this case.


c- Explain your choices of clustering methods and all parameters used. Why did you use this clustering method and parameters? What are its advantages/disadvantages over other potential methods?

```
[6]: from sklearn.metrics import silhouette_score
import part2 as clustering
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans, mean_shift

cluster_data = clustering.get_numerical_features(clean_data)

plt.plot(range(1, 15), clustering.elbow_method(cluster_data), marker='o')
plt.title('Elbow Method')
plt.xlabel('Number of Clusters')
plt.ylabel('Sum of Squared Differences')
plt.show()

kmeans = KMeans(n_clusters=5, init='k-means++', random_state=42)
clusters = kmeans.fit_predict(cluster_data)
silhouette = silhouette_score(cluster_data, clusters)
print(f"Silhouette Score: {silhouette}")
```



Silhouette Score: 0.9135220925159057

KMeans KMeans is widely used clustering algorithms, which aims to partition a dataset into a predefined number of clusters (K). The algorithm minimizes the variance within each cluster. This algorithm was chosen because it is a simple yet very efficient clustering tool, which can handle large datasets. It is specialized on numerical values as well as you can compute the sum of squared distances the best. The disadvantages of KMeans would be the estimation of the number of clusters, which can be solved with the Elbow-Method. Another downside is the assumption of spherical clusters so it doesn't work great for irregular shaped clusters. Also, the algorithm is very sensitive to outliers, which can skew the mean of a cluster.

Here first the table was reduced dimensions with numerical values only and all NAN values were replaced with the mean of the column to not distort the data. After that the first method that was chosen was the KMeans method. For this an estimate of clusters has to be given. So the elbow method was applied. That means K-means was performed, clustering with values from 1-15. For each of the values, we calculate the average distances to the centroid (=center of cluster) across all data points. We want the point where the average distance from the centroid falls suddenly ("elbow"). It was identified to be 5 so it was put in the KMeans function - elbow point can be ambiguous Silhouette Score highest for k=2. The initialization 'k-means++' was chosen due to the advantage of distant centroids rather than completely random initialization. The random state was chosen randomly but provides consistent results even though random initialization.

```
[7]: from sklearn.cluster import DBSCAN

dbscan = DBSCAN(eps=400, min_samples=5)
clusters = dbscan.fit_predict(cluster_data)

if len(set(clusters)) > 1:
    silhouette_avg = silhouette_score(cluster_data, clusters)
    print(f'Silhouette Score: {silhouette_avg}')
else:
    print("Only one cluster or noise, silhouette score cannot be computed.")
```

Silhouette Score: 0.4124997541647688

DBSCAN DBSCAN (Density-Based Spatial Clustering of Applications with Noise) is a clustering algorithm that identifies dense regions in the data and separates out sparse regions (noise). It is particularly effective in scenarios where the data contains clusters of arbitrary shapes and where there is noise. It also has a way of handling outliers - labeling with -1. These points make DBSCAN an interesting choice to compare to KMeans. Unlike methods like K-means, DBSCAN does not require the number of clusters to be specified upfront. The parameters are: `eps` = maximum distance between two samples for them to be considered as in the same neighborhood and `min_samples` = minimum number of points required to form a dense region. A disadvantage of the algorithm is the sensitivity to these parameters, meaning if they are not chosen correctly, clustering might fail. Also, DBSCAN struggles with varying density because it assumes that all clusters have a similar density.

Here we used the same data that only contains numerical features and no NAN values again for a DBSCAN Clustering. The if-else should check if there were meaningful clusters formed and if so calculate the Silhouette Score for the clustering. Starting out with `eps=0.5` (points within distance of 0.5 considered neighbors) and `min_samples=5` (dense region must have ≥ 5 points) all data points were considered 'Noise' (labeled -1). This indicated that the `eps` was too low for the data, which is too sparse for these parameters. So after upping `eps` to 10 the Silhouette Score could be calculated. The results were negative indicating a very bad score. The highest result that could be achieved were with the parameters `eps=400`, `min_samples=5`.

```
[8]: from sklearn.cluster import BisectingKMeans

bisecting_kmeans = BisectingKMeans(n_clusters=5, random_state=42)
clusters = bisecting_kmeans.fit_predict(cluster_data)
silhouette = silhouette_score(cluster_data, clusters)
print(f"Silhouette Score: {silhouette}")
```

Silhouette Score: 0.4703193735377516

Bisecting KMeans Bisecting KMeans is a hierarchical approach of the KMeans algorithm. Its goal is the same but the direction of methodology is the other way around. It doesn't create all centroids at once randomly, but rather picks them progressively based on a previous clustering. The algorithm starts with one cluster and splits until the predefined number of clusters is reached. There are 2 strategies for selecting the cluster to split, the default being "biggest_inertia" will always split cluster with the biggest amount of points assigned to it from all clusters previously

calculated. The second strategy is “largest_cluster” which selects the clusters to split based on having the most points. Both strategies were tried out, and splitting based on biggest inertia resulted in a slightly higher Silhouette Score. It adds the advantage of providing hierarchical structures useful for understanding relationships between clusters, while also reducing sensitivity to initialization by focusing on splitting problematic clusters.

The parameters were chosen to directly compare the hierarchical approach to the regular KMeans algorithm. Although when run with no parameter given for number of clusters a 0.05 higher Silhouette Score was calculated.

Conclusion The decision of which clustering algorithm to choose for this data, relies on the Silhouette Score. It measures how close each point in one cluster is to the points in the neighboring clusters. A score close to +1 indicates that the points are well clustered, a score of 0 means the points are on the border of two clusters, and a negative score indicates that points might have been incorrectly clustered. Looking at the Silhouette Score of KMeans 0.9135220925159057 compared to the Silhouette Score of DBSCAN 0.4124997541647688 or Bisecting Means 0.4703193735377516 the choice seems obvious: KMeans is the best choice for clustering this data.

1.3 Part 3: Dimensionality Reduction and Clustering

Read the documentation of the scikit-learn library for [PCA](#) and [t-SNE](#).

(d) *What are the advantages/disadvantages of these two-dimensionality reduction approaches? Compare the two approaches against each other, theoretically.*

Perform these two-dimensionality reductions on your dataset and visualize the results as a scatterplot. Subsequently, perform a clustering on top of your dimensionality reduction, according to your observations in Part 2 of this exercise. Show your results again in a scatterplot.

[]:

(e) *What are the advantages of using clustering on top of a dimensionality reduction?*

On top of your clustered dimensionality reduced outcome from the previous step, encode in the scatterplot the ‘2345-7_Glucose [Mass/volume] in Serum or Plasma’ variable. For this, you will need to use appropriate visual channels (e.g., opacity, shape, etc.) for the datapoints.

[]:

(f) *Explain the visual variable choices you made in visualizing the glucose variable. Do you have any interesting observations regarding the patterns in the data?*